

Other Systems

Flink & Spark

<https://flink.apache.org/>

<https://spark.apache.org/>

Flink: 1 Slide is Enough

- Unified batch and streaming;
- Event-driven applications (timers, CEP);
- Production-ready SQL support;
- First class state management;
- Fault-tolerance for correctness;
- Connectors for I/O (Kafka, Cassandra);
- Support for multiple deploys (K8s, Yarn, Mesos) and storage systems (HDFS, S3, RocksDB) for snapshots.

Spark

- Fully fledged batch/stream processor as Flink;
- (Now) it also adopts a unified approach;
- Spark had a different vision (it was born with the idea of batch processing), but it is now converging towards Flink ideas (e.g. structured streaming);
- Great support for ML libraries as Flink.

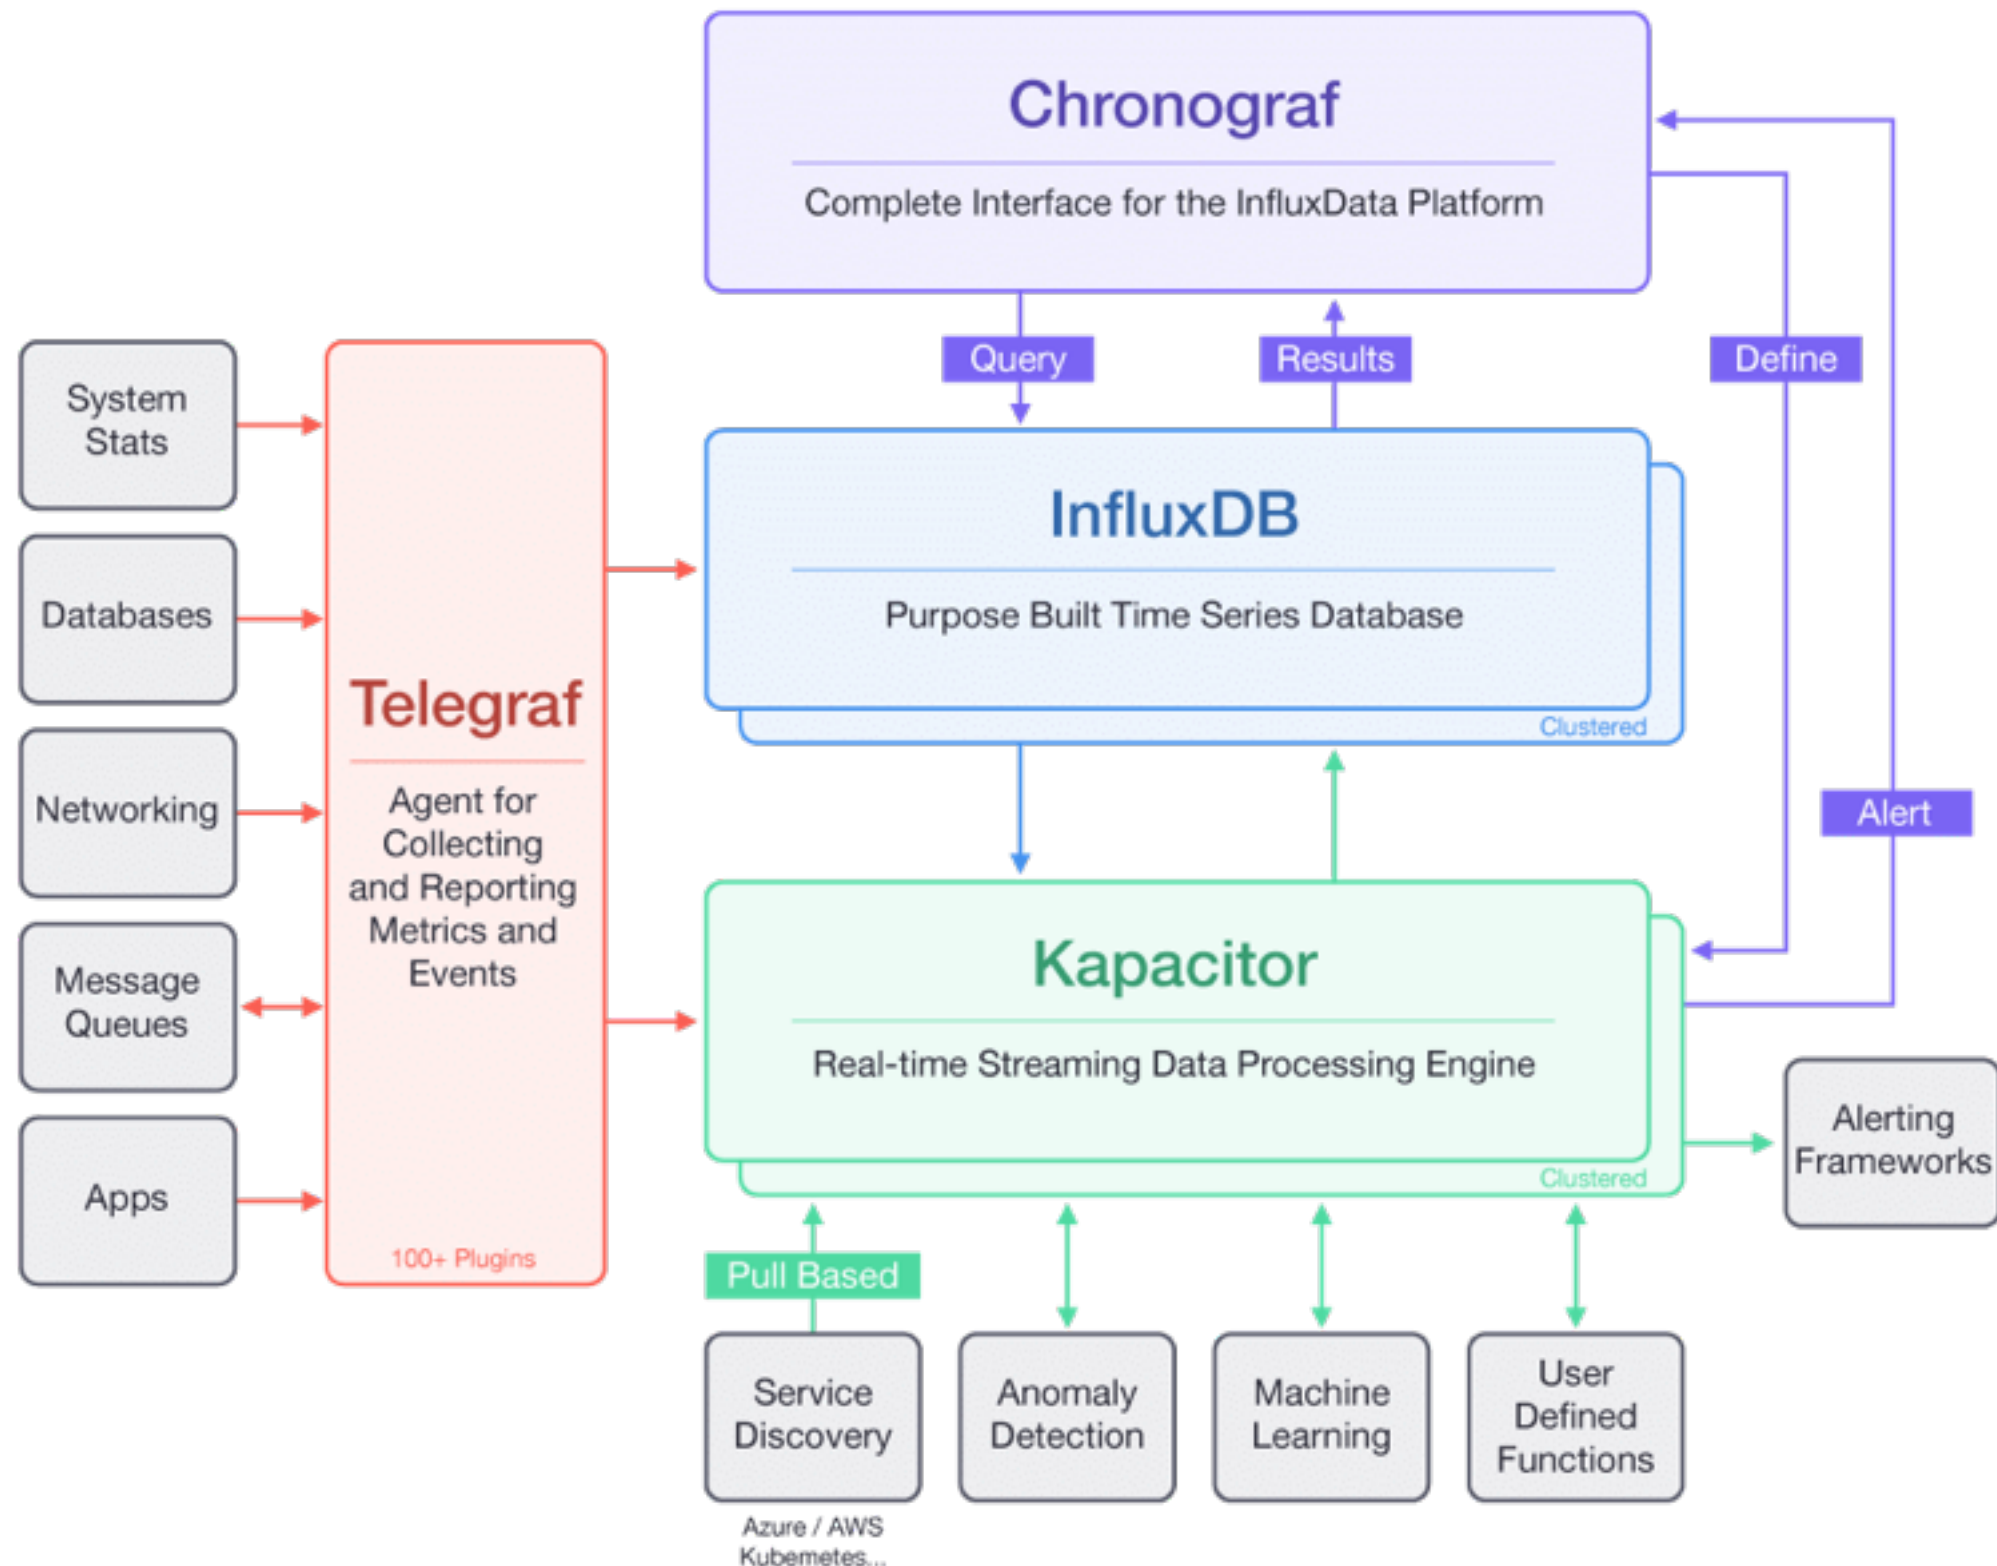
Influx

<https://www.influxdata.com/products/>

The TICK stack

- **T**elegraph: gather data from sources;
- **I**nfluxDB: the core, the Time Series DataBase;
- **C**hronograph: visualize data;
- **K**apacitor: act basing on continuous queries on data;

The TICK stack



TSDB

A Time Series DataBase is optimised for managing timestamped data:

- Timestamp Data compression;
- Optimised for scans over huge *periods*;
- Out-of-the-box data summarization (downsampling);
- Time aware queries.

How to query the Database

- Influx is realizing that static queries, continuous queries and event-processing are very similar;
- They are moving to a SQL-like declarative query language to a unified scripting language: **flux**¹;
- The parser reads the script and generates a plan (as in Flink).

A Flux Query

```
from(db:"foo")
  |> filter(
    fn: (r) => r[ "_measurement" ] == "cpu"
    AND r[ "_field" ] == "usage_system")
  |> range(start:-12h)
  |> window(every:10m, period: 5m)
  |> min()
```

Kapacitor VS CQ

- Influx has a system to perform Continuous Queries
- It provides a SQL-like language for CQs
- Transitioning to TICKscript¹

Kapacitor Script (1/2)

```
dbrrp "telegraf"."autogen"
```

```
stream
```

```
  |from()
```

```
    .measurement('cpu')
```

```
  |window()
```

```
    .period(5m)
```

```
    .every(5m)
```

```
    .align()
```

```
  |mean('usage_idle')
```

```
    .as('usage_idle')
```

Kapacitor Script (2/2)

```
|alert()  
  .crit(lambda: int("usage_idle") < 70)  
  
  // Write it to a file.  
  .log('/tmp/alerts.log')  
  
  // ... or send to REST endpoint  
  .post('https://alerthandler.example.com')  
  
  // Execute script  
  .exec('/bin/custom_alert_handler.sh')
```

Distributed Features (1/2)

- Data is sharded and replicated over nodes;
- *Shard groups* are based on *retention policy* (the time that data loses importance);
- On write, data is sent to the right shard within the shard group by using the *tag set* of the record (like indexing in normal DBs, or `keyBy` in Flink);
- If your retention policy is 1 day, then, for any given tag set, all data on any given day lies in a single shard;
- Queries that span multiple days must be *reduced* by the coordinator.

Distributed Features (2/2)

- Different levels of consistency for replication:
 - any: the write has been logged (*hinted handoff*)
 - one: the write is on at least one replica
 - quorum: the write is on the majority of replicas
 - all: synchronous replication
- Fault-tolerance: hinted handoff is a persisted log (for every node to every node) that is replayed in case of failure.

A Final Note

- Single machine is open-source;
- Distributed cluster requires payment;
- The database is distributed, the computation is not;
- Every query is executed by a single node, if it fails, then it fails as a whole.

Kafka

<https://kafka.apache.org/>


About Kafka

- **Kafka** *is a*¹ distributed streaming platform;
- Its core is a highly scalable and fault tolerant log store;
- **Kafka Streams** is a library to write streaming pipelines consume kafka logs and produce kafka logs.

Kafka Concepts (1/3)

The topic: a partitioned stream

Anatomy of a Topic



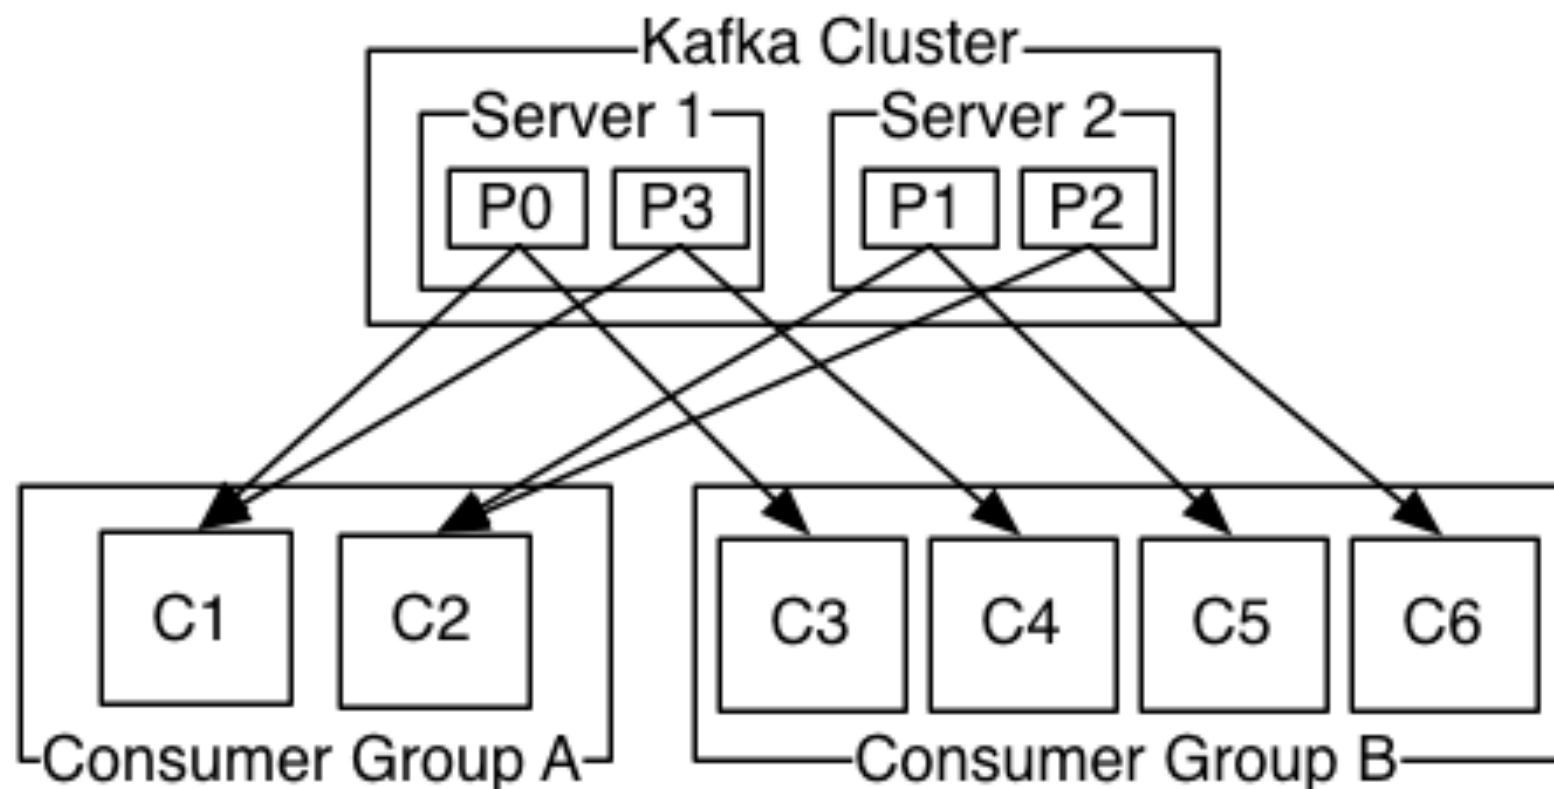
- Every partition is an append-only log;
- Every record is assigned an *offset* in the partition.

Kafka Concepts (2/3)

- Partitions allow for scale;
- Every partition can be replicated for fault-tolerance (master-slave replication);
- The producers of data pushes data to some partition basing on some policy (e.g. it could use a key).

Kafka Concepts (3/3)


- Consumers are organized in groups;
- Every consumer in a consumer group is assigned to a partition;
- Consumers can join/leave a group —> partition assignment is handled transparently.



Kafka Streams

- A simple library to write stream processing jobs;
- There is NO runtime —> no schedulers, no taskmanagers, no jobmanagers;
- Kafka Streams is built on top of the consumer group concept;
- There is a stream processing *task* for each partition in the input topic(s);
- Each consumer in the consumer group runs some of the tasks and produces to other topics;

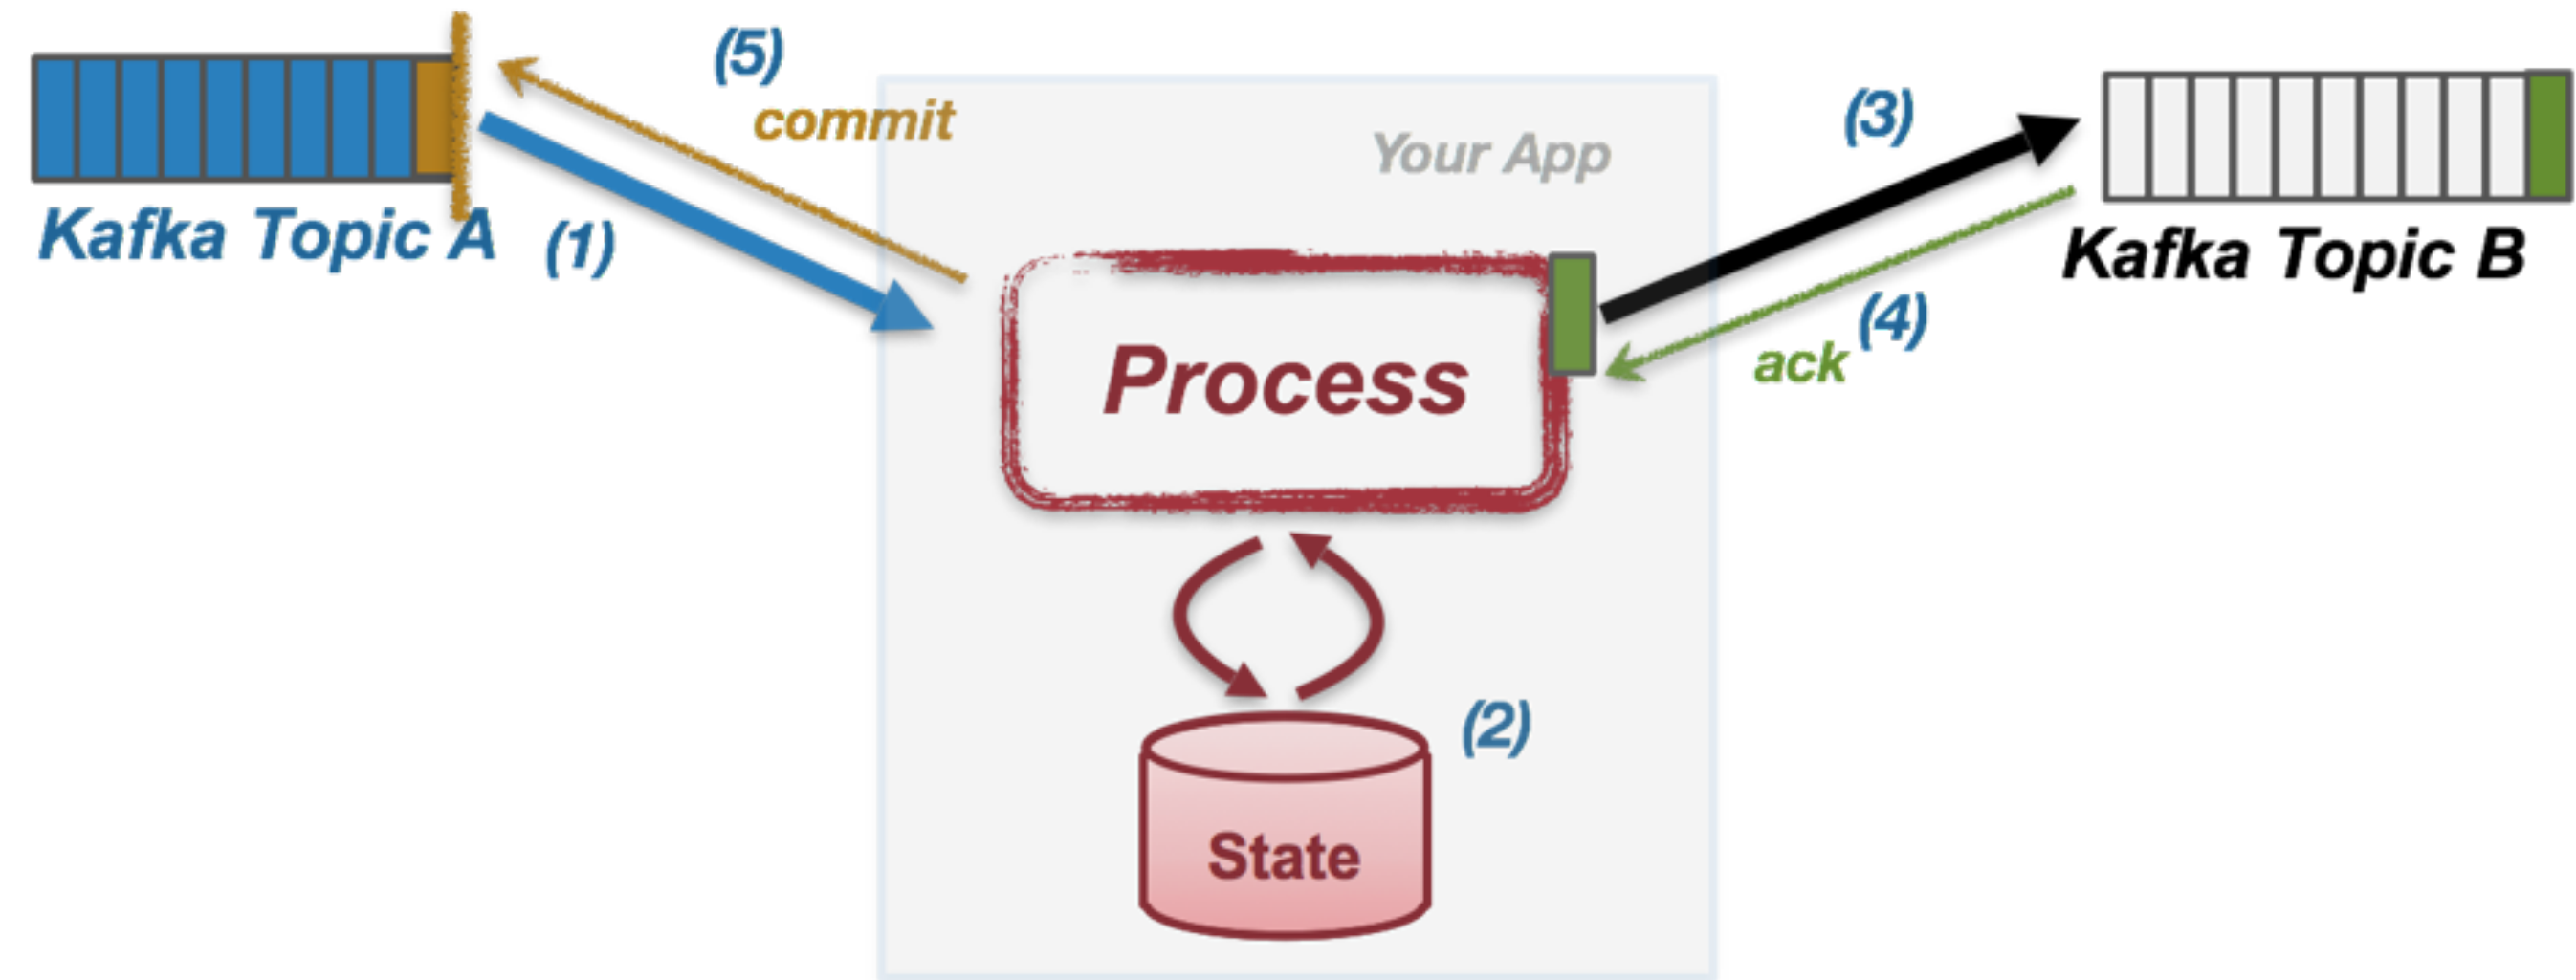
Kafka Streams



Kafka Streams - Failure (1/2)

- Local state changes are persisted to dedicated kafka topics too!
- If some consumer fails, its tasks are assigned to another one that can replay the state changes;
- Kafka guarantees exactly-once processing guarantees among multiple kafka topics by storing offsets and committing records once they are produced towards the next kafka topic;
- NOTE that re-partitioning causes to write to a kafka topic (again).

Kafka Streams - Failure (2/2)



Differences wrt Flink? Opinions?

KSQL (1/2)

SQL-like language for streaming

```
CREATE STREAM pageviews_enriched  
AS |  
  SELECT pv.viewtime, |  
         pv.userid AS userid, |  
         ...  
         u.contact_info |  
  FROM page_views pv |  
  LEFT JOIN users u ON pv.userid =  
  users.userid;
```

KSQL (2/2)

SQL-like language for streaming

```
CREATE TABLE  
windowed AS \  
SELECT regionid, count(*) \  
FROM pageviews_enriched \  
WINDOW TUMBLING (SIZE 1 MINUTE) \  
GROUP BY regionid;
```

Connectors

- Kafka Connect is a separate project¹ that aims at integrating Kafka broker with other systems;
- Source and sink connectors;
- Some connectors are developed by Confluent (Kafka's company), some are certified, some are from the community;
- To name a few:
 - JDBC, Elasticsearch, JMS, HDFS;
 - VoltDB, MongoDB, Redis, RabbitMQ.