

# Streams & Tables

Most of the images used in this section are courtesy of the [Flink Documentation](#)

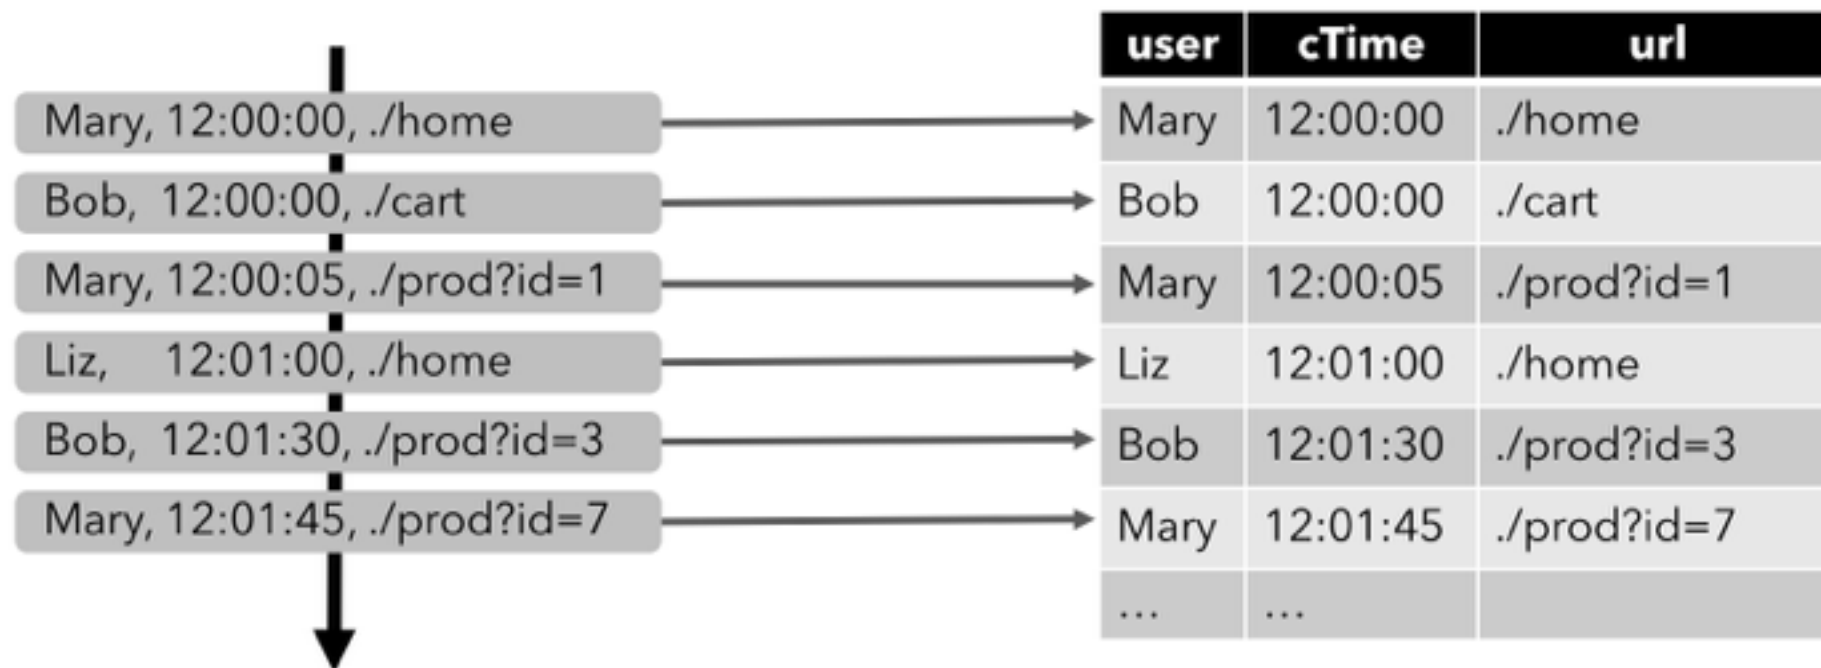
# Dynamic Tables



- Tables changing over time
- $S2R \rightarrow R2R \rightarrow R2S$
- How to pass from one to the other?

# S2R

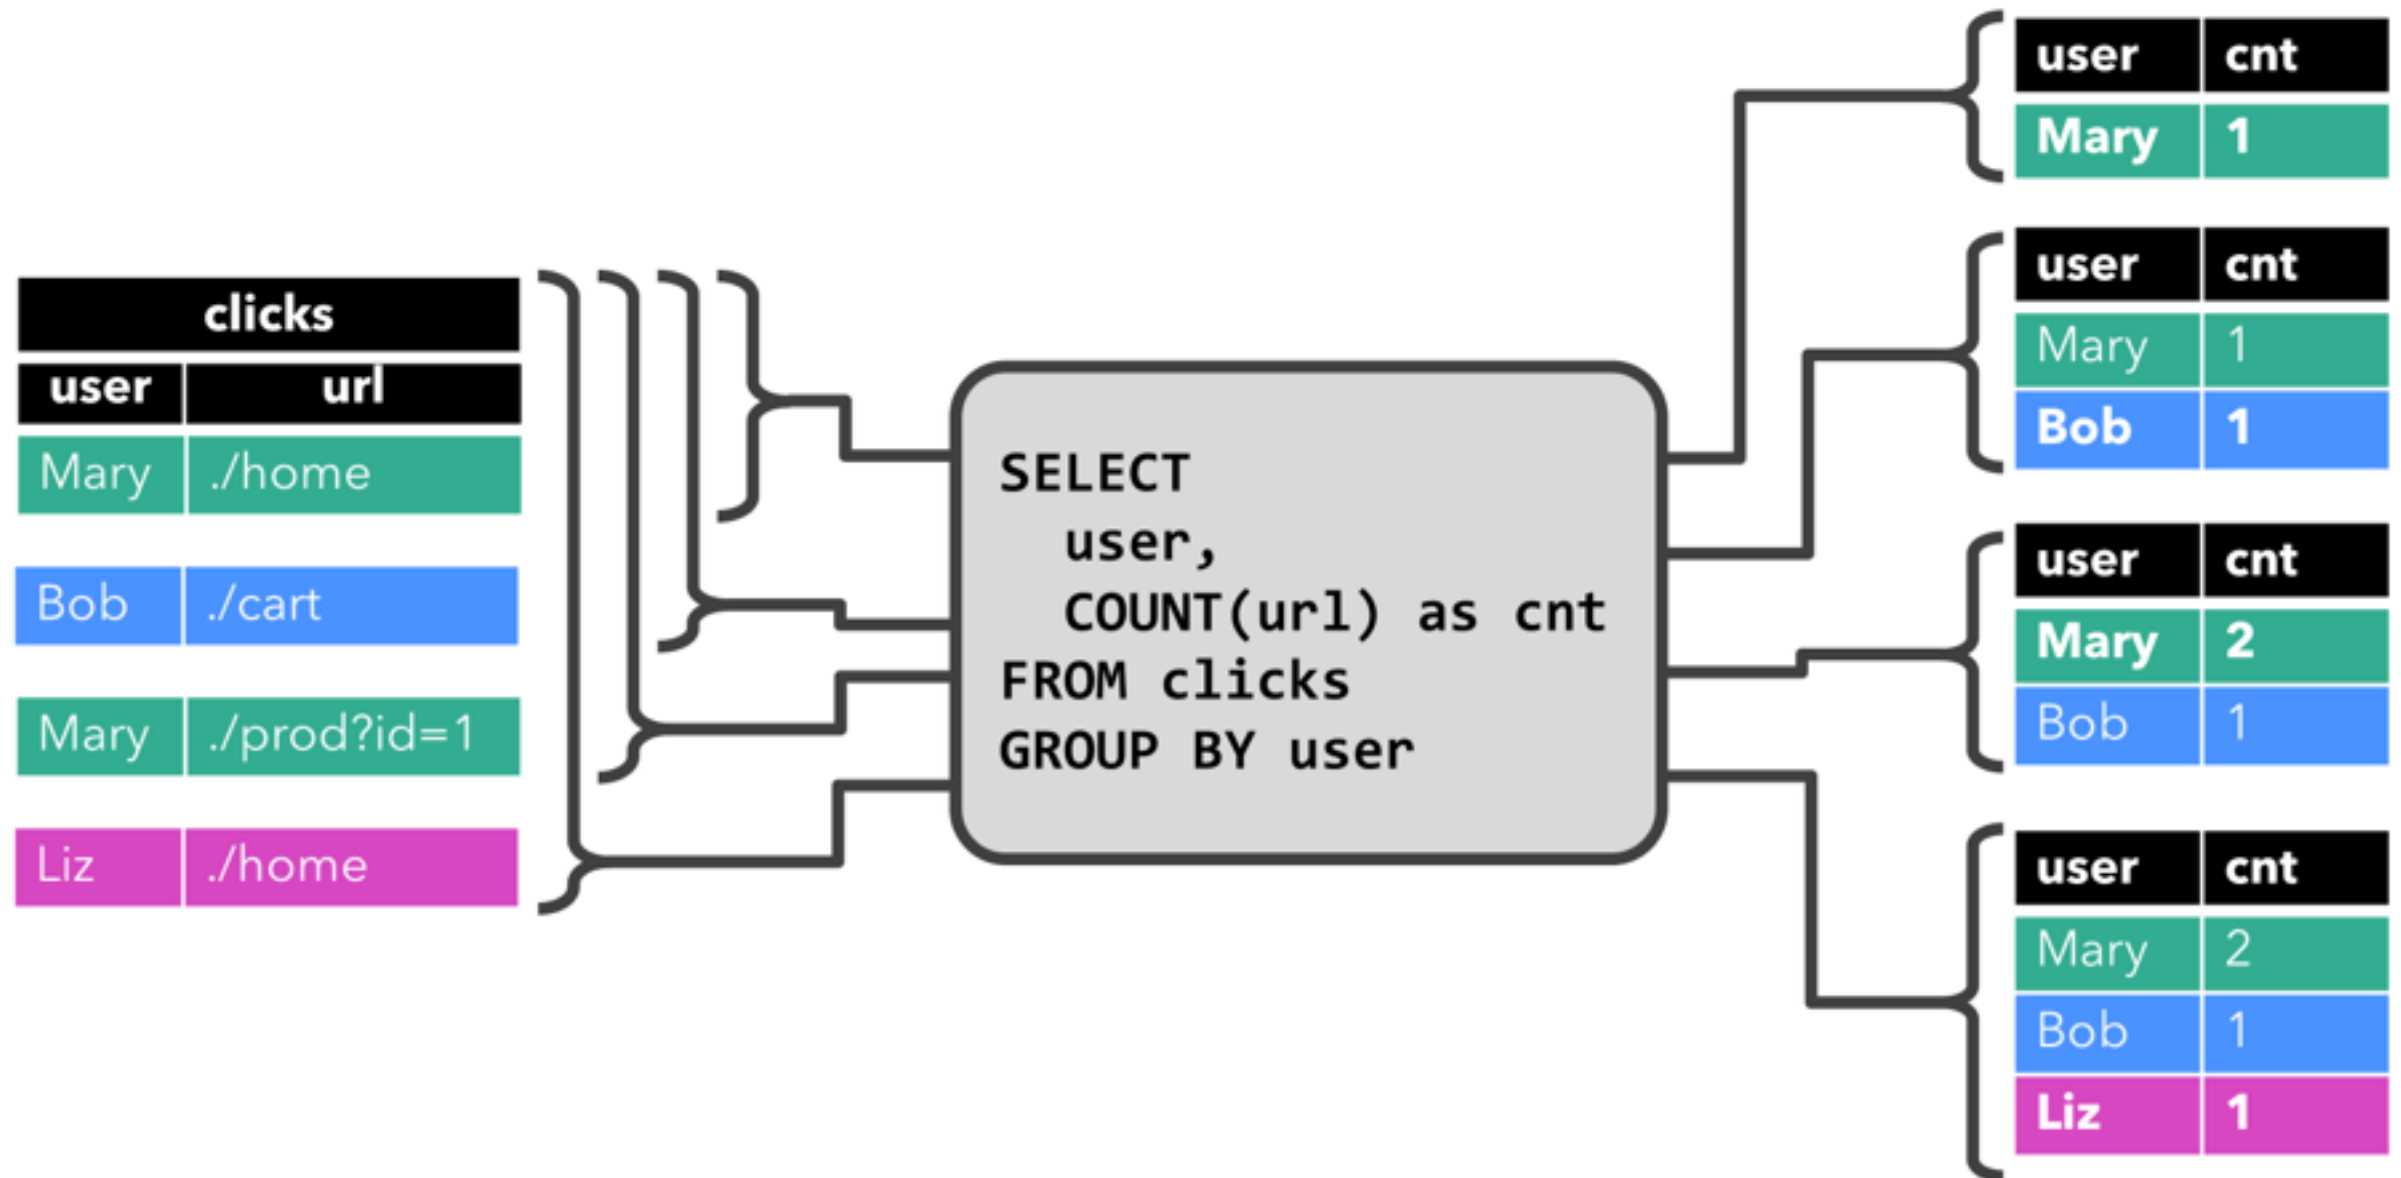
**Stream to Relation** operator converts a stream into a relation (dynamic table) by appending every record to the table



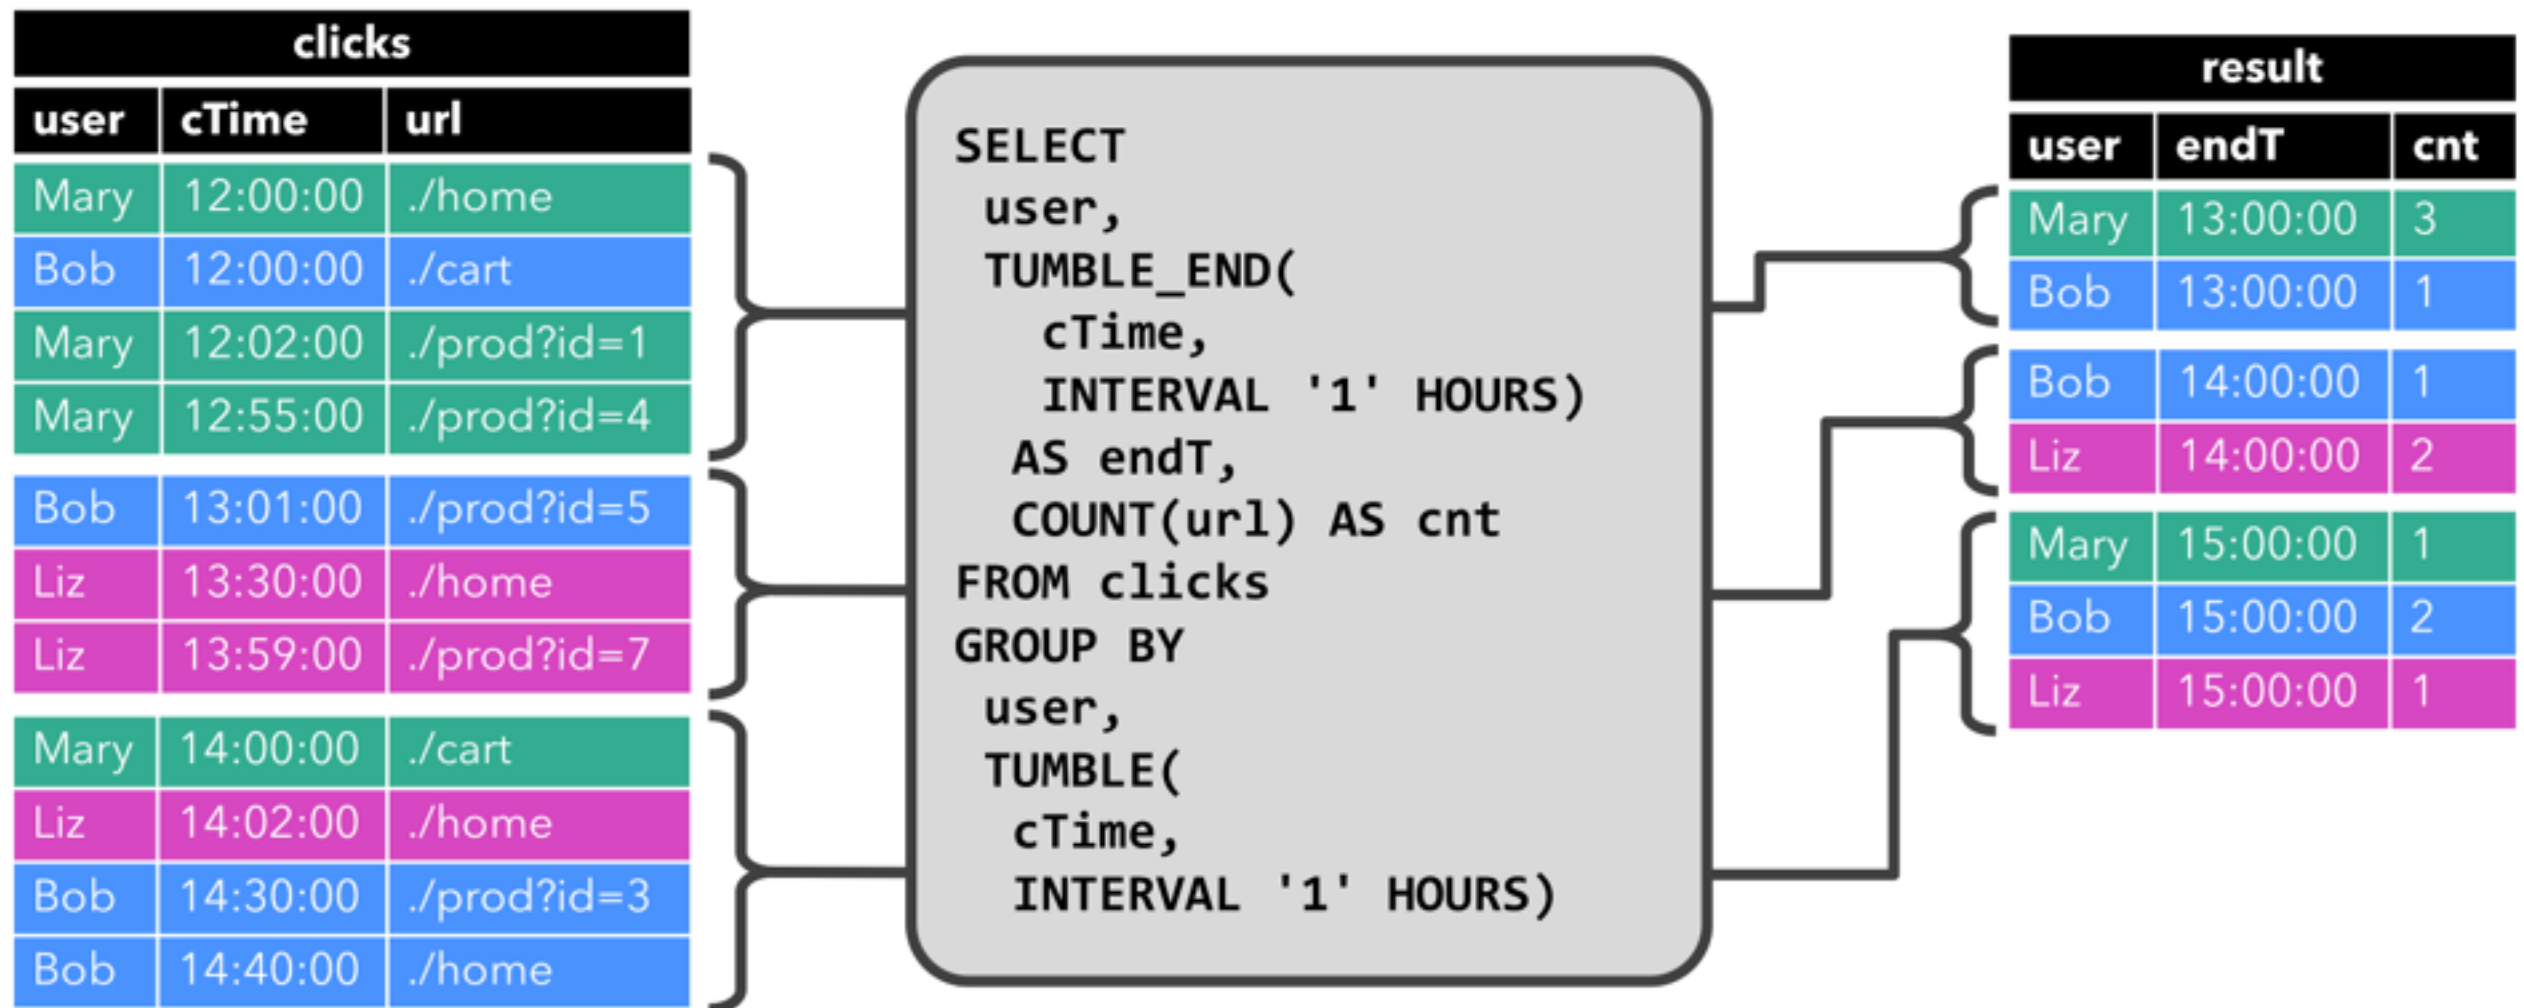
# R2R

- **Relation to Relation** operators are the classic SQL operators that given a table provide a table as result;
- If the input table is dynamic (because it records a stream) the SQL query becomes a continuous query;
- SQL queries on dynamic tables can involve time-dependent transformations (windows).

# R2R - Continuous Query (1)

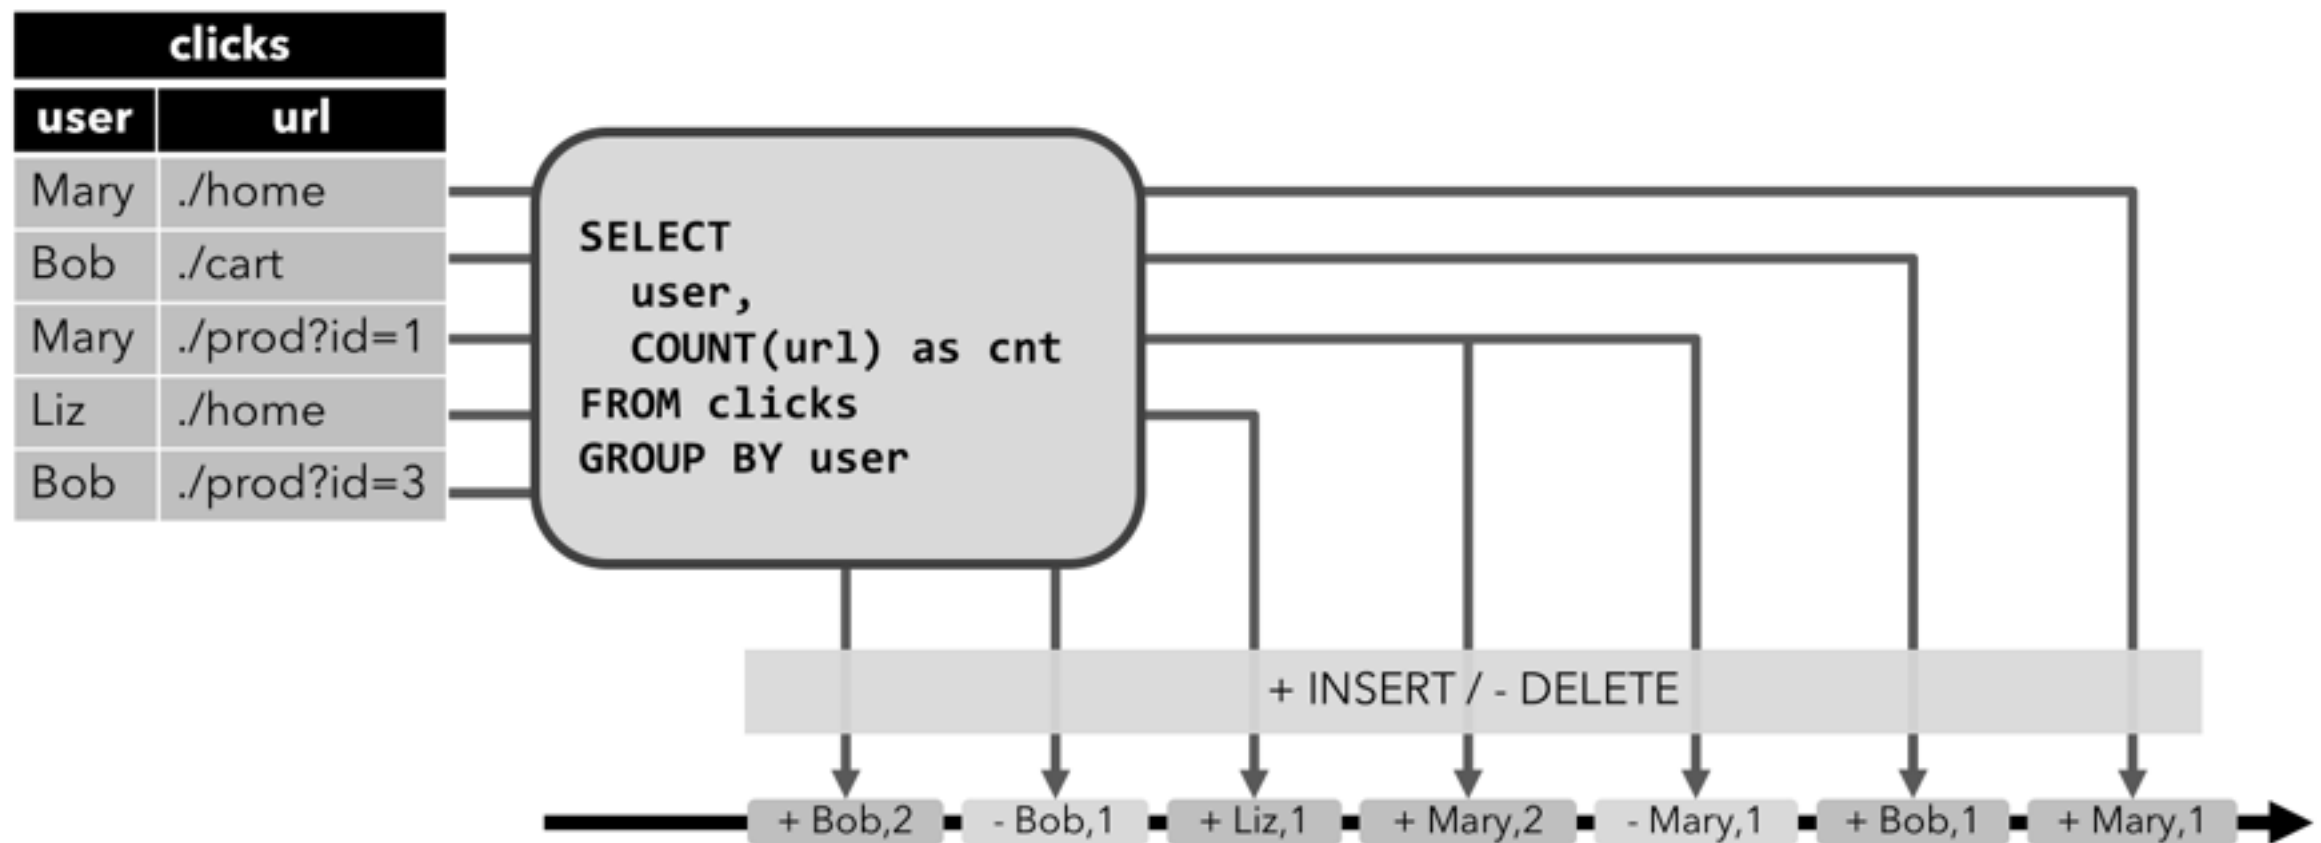


# R2R - Continuous Query (2)



# R2S

**Relation to Stream** convert a dynamic table to a stream as a log of updates (insert/update)



# API

```
// for batch programs use ExecutionEnvironment instead of StreamExecutionEnvironment
StreamExecutionEnvironment env = StreamExecutionEnvironment.getExecutionEnvironment();

// create a TableEnvironment
// for batch programs use BatchTableEnvironment instead of StreamTableEnvironment
StreamTableEnvironment tableEnv = TableEnvironment.getTableEnvironment(env);

// register a Table
tableEnv.registerTable("table1", ...)           // or
tableEnv.registerTableSource("table2", ...);    // or
tableEnv.registerExternalCatalog("extCat", ...);

// create a Table from a Table API query
Table tapiResult = tableEnv.scan("table1").select(...);
// create a Table from a SQL query
Table sqlResult  = tableEnv.sql("SELECT ... FROM table2 ... ");

// emit a Table API result Table to a TableSink, same for SQL result
tapiResult.writeToSink(...);

// execute
env.execute();
```




# CQL

- Is the **Continuous Query Language** created by Jennifer Widom in 2003<sup>1</sup>;
- I intentionally used **S2R**, **R2R**, and **R2S** because it is the notation that Widom used;
- The major difference between CQL and Flink Dynamic Tables lies in S2R: in CQL the stream to relation operator is the window;
- R2R is simply static SQL.


# The CEP Library

# Patterns



# Conditions

Shape is Square




```
p.subtype(Square.class)
```

and is blue




```
.where(  
    s -> s.color() == Color.BLUE  
)
```

and is transparent





```
.where(...) // AND
```

or is green



```
.or(...) // OR
```

# Iterative Conditions



“Accept this square  
if the average  
surface **of the last  
accepted shapes** is  
less than 10”


# Quantifiers



singleton (a)



... looping (a+)



... optional (a?)

`p.times(n)`   `p.oneOrMore()`   `p.optional()`

# Contiguity (1/2)

- **Strict Contiguity:** Expects all matching events to appear strictly one after the other, without any non-matching events in-between.
- **Relaxed Contiguity:** Ignores non-matching events appearing in-between the matching ones.
- **Non-Deterministic Relaxed Contiguity:** Further relaxes contiguity, allowing additional matches that ignore some matching events.

# Contiguity (2/2)

Pattern: “ $a + b$ ”, Input: “ $a1 - c - a2 - b$ ”


- **Strict Contiguity:**  $\{a2 - b\}$  – the  $c$  after  $a1$  causes  $a1$  to be discarded.
- **Relaxed Contiguity:**  $\{a1 - b\}$  and  $\{a1 - a2 - b\}$  –  $c$  is ignored.
- **Non-Deterministic Relaxed Contiguity:**  $\{a1 - b\}$ ,  $\{a2 - b\}$ , and  $\{a1 - a2 - b\}$ .




# Contiguity on Loops




`square.oneOrMore().?.circle`



strict contiguity  
`.consecutive()`




relaxed contiguity




non-deterministic  
contiguity  
`.allowCombinations()`

# Combining Patterns


`square.oneOrMore().?.?.circle`



strict contiguity  
`.next()`



relaxed contiguity  
`.followedBy()`



non-deterministic  
contiguity  
`.followedByAny()`

NOT

# Temporal Constraints



```
square  
  .followedBy(circle)  
  .within(Time.seconds(10))
```

# Example

```
StreamExecutionEnvironment env = ...
env.setStreamTimeCharacteristic(TimeCharacteristic.EventTime);

DataStream<Event> input = ...

DataStream<Event> partitionedInput = input.keyBy(new KeySelector<Event, Integer>() {
    @Override
    public Integer getKey(Event value) throws Exception {
        return value.getId();
    }
});

Pattern<Event, ?> pattern = Pattern.<Event>begin("start")
    .next("middle").where(new SimpleCondition<Event>() {
        @Override
        public boolean filter(Event value) throws Exception {
            return value.getName().equals("error");
        }
    }).followedBy("end").where(new SimpleCondition<Event>() {
        @Override
        public boolean filter(Event value) throws Exception {
            return value.getName().equals("critical");
        }
    }).within(Time.seconds(10));

PatternStream<Event> patternStream = CEP.pattern(partitionedInput, pattern);

DataStream<Alert> alerts = patternStream.select(new PatternSelectFunction<Event, Alert>() {
    @Override
    public Alert select(Map<String, List<Event>> pattern) throws Exception {
        return createAlert(pattern);
    }
});
```

# References

- [More on tables and streams](#)
- [CEP Library](#)